

Package ‘LinkOrgs’

January 28, 2024

Title LinkOrgs: Algorithms for Organizational Record Linkage

Version 0.01

Description An R package for organizational records using the algorithms of Jerzak & Libgobler (2023+). The linkage is done based on organizational names and using half a billion open collaborated records on those names from LinkedIn users. It also contains functions implementing string matching performance metrics, as well as a fast, parallized version of fuzzy string matching.

Depends R (>= 3.3.3)

License

Creative Commons Attribution-Noncommercial-No Derivative Works 4.0, for academic use only.

Encoding UTF-8

LazyData true

Maintainer 'Connor Jerzak' <connor.jerzak@gmail.com>

Imports

data.table,plyr,Rfast,stringdist,doMC,parallel,glmnet,parallel,stringr,dplyr,fastmatch,reticulate

RoxygenNote 7.2.3

R topics documented:

BuildBackend	1
BuildTransfer	2
GetCalibratedDistThres	3
LinkOrgs	4
pDistMatch	7
pFuzzyMatch	8
url2dt	11
Index	12

BuildBackend	<i>Build the environment for LinkOrgs machine learning models. Builds a conda environment in which jax, optax, equinox, and jmp are installed.</i>
--------------	--

Description

Build the environment for LinkOrgs machine learning models. Builds a conda environment in which jax, optax, equinox, and jmp are installed.

Usage

```
BuildBackend(conda_env = "LinkOrgsEnv", conda = "auto")
```

Arguments

- conda_env (default = "LinkOrgsEnv") Name of the conda environment in which to place the backends.
- conda (default = auto) The path to a conda executable. Using "auto" allows reticulate to attempt to automatically find an appropriate conda binary.

Value

Builds the computational environment for LinkOrgs. This function requires an Internet connection. You may find out a list of conda Python paths via: `system("which python")`

Examples

```
# For a tutorial, see
# github.com/cjerzak/linkorgs-software/
```

BuildTransfer	<i>A primarily internal function which builds the organizational record linkage models used in Libgober and Jerzak (2023+).</i>
---------------	---

Description

A primarily internal function which builds the organizational record linkage models used in Libgober and Jerzak (2023+).

Usage

```
BuildTransfer()
```

GetCalibratedDistThres

GetCalibratedDistThres

Description

Performs parallelized fuzzy matching of strings based on the string distance measure specified in DistanceMeasure. Matching is parallelized using all available CPU cores to increase execution speed.

Usage

```
GetCalibratedDistThres(x, y, AveMatchNumberPerAlias = 5L)
```

Arguments

x, y	data frames to be merged
by, by.x, by.y	specifications of the columns used for merging. We follow the general syntax of <code>base::merge</code> ; see <code>?base::merge</code> for more details.
...	For additional options, see “Details”.

Details

pFuzzyMatch can automatically process the by text for each dataset. Users may specify the following options:

- Set DistanceMeasure to control algorithm for computing pairwise string distances. Options include "osa", "jaccard", "jw". See `?stringdist::stringdist` for all options. (Default is "jaccard")
- Set MaxDist to control the maximum allowed distance between two matched strings
- Set AveMatchNumberPerAlias to control the maximum allowed distance between two matched strings. Takes priority over MaxDist if both specified.
- Set qgram to control the character-level q-grams used in the distance measure. (Default is 2)
- Set RemoveCommonWords to TRUE to remove common words (those appearing in > 10% of aliases). (Default is FALSE)
- Set NormalizeSpaces to TRUE to remove hanging whitespaces. (Default is TRUE)
- Set RemovePunctuation to TRUE to remove punctuation. (Default is TRUE)
- Set ToLower to TRUE to ignore case. (Default is TRUE)

Value

z The merged data frame.

Examples

```
#Create synthetic data
x_orenames <- c("apple","oracle","enron inc.,"mcdonalds corporation")
y_orenames <- c("apple corp","oracle inc","enron","mcdonalds co")
x <- data.frame("orenames_x"=x_orenames)
y <- data.frame("orenames_y"=y_orenames)
z <- data.frame("orenames_x"=x_orenames[1:2], "orenames_y"=y_orenames[1:2])
z_true <- data.frame("orenames_x"=x_orenames, "orenames_y"=y_orenames)

# Perform merge
linkedOrgs_fuzzy <- pFuzzyMatch(x = x,
                                y = y,
                                by.x = "orenames_x",
                                by.y = "orenames_y")
```

LinkOrgs

LinkOrgs

Description

Implements the organizational record linkage algorithms of Libgober and Jerzak (2023+) using half-a-billion open-collaborated records.

Implements the organizational record linkage algorithms of Libgober and Jerzak (2023+) using half-a-billion open-collaborated records.

Usage

```
LinkOrgs(
  x,
  y,
  by = NULL,
  by.x = NULL,
  by.y = NULL,
  algorithm = "ml",
  conda_env = NULL,
  conda_env_required = F,
  ReturnDiagnostics = F,
  ReturnProgress = T,
  ToLower = T,
  NormalizeSpaces = T,
  RemovePunctuation = T,
  MaxDist = NULL,
  MaxDist_network = NULL,
  AveMatchNumberPerAlias = 10,
  AveMatchNumberPerAlias_network = 2,
  DistanceMeasure = "jaccard",
  qgram = 2,
  ml_version = "v0",
  openBrowser = F,
  ReturnDecomposition = F
```

```

)

LinkOrgs(
  x,
  y,
  by = NULL,
  by.x = NULL,
  by.y = NULL,
  algorithm = "ml",
  conda_env = NULL,
  conda_env_required = F,
  ReturnDiagnostics = F,
  ReturnProgress = T,
  ToLower = T,
  NormalizeSpaces = T,
  RemovePunctuation = T,
  MaxDist = NULL,
  MaxDist_network = NULL,
  AveMatchNumberPerAlias = 10,
  AveMatchNumberPerAlias_network = 2,
  DistanceMeasure = "jaccard",
  qgram = 2,
  ml_version = "v0",
  openBrowser = F,
  ReturnDecomposition = F
)

```

Arguments

<code>x, y</code>	data frames to be merged
<code>by, by.x, by.y</code>	character vector(s) that specify the column names used for merging data frames <code>x</code> and <code>y</code> . The merging variables should be organizational names. See <code>?base::merge</code> for more details regarding syntax.
<code>algorithm</code>	character; specifies which algorithm described in Libgober and Jerzak (2023+) should be used. Options are "markov", "bipartite", "ml", and "transfer". Default is "ml", which uses a machine-learning approach using Transformer netes and 9 million parameters to predict match probabilities using half a billion open-collaborated records as training data.
<code>conda_env</code>	character string; specifies a conda environment where JAX and related packages have been installed (see <code>?LinkOrgs::BuildBackend</code>). Used only when <code>algorithm='ml'</code> or <code>DistanceMeasure='ml'</code> .
<code>conda_env_required</code>	Boolean; specifies whether conda environment is required.
<code>ReturnDiagnostics</code>	logical; specifies whether various match-level diagnostics should be returned in the merged data frame.
<code>ml_version</code>	character; specifies which version of the ML algorithm should be used. Options are of the form "v1", "v2", "v3" Highest version currently supported is "v1" (11M parameters).
...	For additional specification options, see "Details".

Details

LinkOrgs automatically processes the name text for each dataset (specified by `by` or `by.x`, and `by.y`). Users may specify the following options:

- Set `DistanceMeasure` to control algorithm for computing pairwise string distances. Options include "osa", "jaccard", "jw". See `?stringdist::stringdist` for all options. Default is "jaccard". To use the combined machine learning and network methods, set `algorithm` to "bipartite" or "markov", and `DistanceMeasure` to "ml".
- Set `MaxDist` to control the maximum allowed distance between two matched strings
- Set `MaxDist_network` to control the maximum allowed distance between two matched strings in the integration with the LinkedIn network representation.
- Set `AveMatchNumberPerAlias` to control the maximum allowed distance between two matched strings. Takes priority over `MaxDist` if both specified.
- Set `AveMatchNumberPerAlias_network` to control the maximum allowed distance between two matched strings in the integration with the LinkedIn network representation. Takes priority over `MaxDist_network` if both specified.
- Set `qgram` to control the character-level q-grams used in the distance measure. Default is 2.
- Set `RemoveCommonWords` to TRUE to remove common words (those appearing in > 10% of aliases). Default is FALSE.
- Set `NormalizeSpaces` to TRUE to remove hanging whitespaces. Default is TRUE.
- Set `RemovePunctuation` to TRUE to remove punctuation. Default is TRUE.
- Set `ToLower` to TRUE to ignore case. Default is TRUE.

LinkOrgs automatically processes the name text for each dataset (specified by `by` or `by.x`, and `by.y`). Users may specify the following options:

- Set `DistanceMeasure` to control algorithm for computing pairwise string distances. Options include "osa", "jaccard", "jw". See `?stringdist::stringdist` for all options. Default is "jaccard". To use the combined machine learning and network methods, set `algorithm` to "bipartite" or "markov", and `DistanceMeasure` to "ml".
- Set `MaxDist` to control the maximum allowed distance between two matched strings
- Set `MaxDist_network` to control the maximum allowed distance between two matched strings in the integration with the LinkedIn network representation.
- Set `AveMatchNumberPerAlias` to control the maximum allowed distance between two matched strings. Takes priority over `MaxDist` if both specified.
- Set `AveMatchNumberPerAlias_network` to control the maximum allowed distance between two matched strings in the integration with the LinkedIn network representation. Takes priority over `MaxDist_network` if both specified.
- Set `qgram` to control the character-level q-grams used in the distance measure. Default is 2.
- Set `RemoveCommonWords` to TRUE to remove common words (those appearing in > 10% of aliases). Default is FALSE.
- Set `NormalizeSpaces` to TRUE to remove hanging whitespaces. Default is TRUE.
- Set `RemovePunctuation` to TRUE to remove punctuation. Default is TRUE.
- Set `ToLower` to TRUE to ignore case. Default is TRUE.

Value

- z The merged data frame.
- z The merged data frame.

Examples

```
#Create synthetic data
x_orenames <- c("apple","oracle","enron inc.,"mcdonalds corporation")
y_orenames <- c("apple corp","oracle inc","enron","mcdonalds co")
x <- data.frame("orenames_x"=x_orenames)
y <- data.frame("orenames_y"=y_orenames)

# Perform merge
linkedOrgs <- LinkOrgs(x = x,
                       y = y,
                       by.x = "orenames_x",
                       by.y = "orenames_y",
                       MaxDist = 0.6)

print( linkedOrgs )

#Create synthetic data
x_orenames <- c("apple","oracle","enron inc.,"mcdonalds corporation")
y_orenames <- c("apple corp","oracle inc","enron","mcdonalds co")
x <- data.frame("orenames_x"=x_orenames)
y <- data.frame("orenames_y"=y_orenames)

# Perform merge
linkedOrgs <- LinkOrgs(x = x,
                       y = y,
                       by.x = "orenames_x",
                       by.y = "orenames_y",
                       MaxDist = 0.6)

print( linkedOrgs )
```

pDistMatch

pDistMatch

Description

Performs parallelized distance computation strings based on the string distance measure specified in DistanceMeasure. Matching is parallelized using all available CPU cores to increase execution speed.

Usage

```
pDistMatch(embed1, embed2, AcceptThreshold = NULL)
```

Arguments

x, y	data frames to be merged
by, by.x, by.y	specifications of the columns used for merging. We follow the general syntax of <code>base::merge</code> ; see <code>?base::merge</code> for more details.
...	For additional options, see “Details”.

Details

pDistMatch can automatically process the by text for each dataset. Users may specify the following options:

- Set DistanceMeasure to control algorithm for computing pairwise string distances. Options include "osa", "jaccard", "jw". See ?stringdist::stringdist for all options. (Default is "jaccard")

Value

z The merged data frame.

Examples

```
#Create synthetic data
x_orenames <- c("apple", "oracle", "enron inc.", "mcdonalds corporation")
y_orenames <- c("apple corp", "oracle inc", "enron", "mcdonalds co")
x <- data.frame("orenames_x"=x_orenames)
y <- data.frame("orenames_y"=y_orenames)
z <- data.frame("orenames_x"=x_orenames[1:2], "orenames_y"=y_orenames[1:2])
z_true <- data.frame("orenames_x"=x_orenames, "orenames_y"=y_orenames)

# Perform merge
linkedOrgs_fuzzy <- pFuzzyMatch(x = x,
                                y = y,
                                by.x = "orenames_x",
                                by.y = "orenames_y")
```

pFuzzyMatch

pFuzzyMatch

Description

Performs fast fuzzy matching of strings based on the string distance measure specified in DistanceMeasure. Matching is parallelized using all available CPU cores to increase execution speed.

Performs parallelized fuzzy matching of strings based on the string distance measure specified in DistanceMeasure. Matching is parallelized using all available CPU cores to increase execution speed.

Usage

```
pFuzzyMatch(
  x,
  y,
  by = NULL,
  by.x = NULL,
  by.y = NULL,
  return_stringdist = T,
  onlyUFT = T,
  qgram = 2,
  DistanceMeasure = "jaccard",
```



```

    MaxDist = 0.2,
    AverageReference = NULL,
    AveMatchNumberPerAlias = NULL,
    openBrowser = F,
    ReturnProgress = T,
    ReturnMaxDistThreshold = F
  )

pFuzzyMatch(
  x,
  y,
  by = NULL,
  by.x = NULL,
  by.y = NULL,
  return_stringdist = T,
  onlyUTF = T,
  qgram = 2,
  DistanceMeasure = "jaccard",
  MaxDist = 0.2,
  AverageReference = NULL,
  AveMatchNumberPerAlias = NULL,
  openBrowser = F,
  ReturnProgress = T,
  ReturnMaxDistThreshold = F
)

```

Arguments

<code>x, y</code>	data frames to be merged
<code>by, by.x, by.y</code>	specifications of the columns used for merging. We follow the general syntax of <code>base::merge</code> ; see <code>?base::merge</code> for more details.
<code>...</code>	For additional options, see “Details”.

Details

pFuzzyMatch can automatically process the `by` text for each dataset. Users may specify the following options:

- Set `DistanceMeasure` to control algorithm for computing pairwise string distances. Options include "osa", "jaccard", "jw". See `?stringdist::stringdist` for all options. (Default is "jaccard")
- Set `MaxDist` to control the maximum allowed distance between two matched strings
- Set `AveMatchNumberPerAlias` to control the maximum allowed distance between two matched strings. Takes priority over `MaxDist` if both specified.
- Set `qgram` to control the character-level q-grams used in the distance measure. (Default is 2)
- Set `RemoveCommonWords` to TRUE to remove common words (those appearing in > 10% of aliases). (Default is FALSE)
- Set `NormalizeSpaces` to TRUE to remove hanging whitespaces. (Default is TRUE)
- Set `RemovePunctuation` to TRUE to remove punctuation. (Default is TRUE)
- Set `ToLower` to TRUE to ignore case. (Default is TRUE)

pFuzzyMatch can automatically process the by text for each dataset. Users may specify the following options:

- Set DistanceMeasure to control algorithm for computing pairwise string distances. Options include "osa", "jaccard", "jw". See ?stringdist::stringdist for all options. (Default is "jaccard")
- Set MaxDist to control the maximum allowed distance between two matched strings
- Set AveMatchNumberPerAlias to control the maximum allowed distance between two matched strings. Takes priority over MaxDist if both specified.
- Set qgram to control the character-level q-grams used in the distance measure. (Default is 2)
- Set RemoveCommonWords to TRUE to remove common words (those appearing in > 10% of aliases). (Default is FALSE)
- Set NormalizeSpaces to TRUE to remove hanging whitespaces. (Default is TRUE)
- Set RemovePunctuation to TRUE to remove punctuation. (Default is TRUE)
- Set ToLower to TRUE to ignore case. (Default is TRUE)

Value

z The merged data frame.

z The merged data frame.

Examples

```
#Create synthetic data
x_orenames <- c("apple","oracle","enron inc.,"mcdonalds corporation")
y_orenames <- c("apple corp","oracle inc","enron","mcdonalds co")
x <- data.frame("orenames_x"=x_orenames)
y <- data.frame("orenames_y"=y_orenames)
z <- data.frame("orenames_x"=x_orenames[1:2], "orenames_y"=y_orenames[1:2])
z_true <- data.frame("orenames_x"=x_orenames, "orenames_y"=y_orenames)

# Perform merge
linkedOrgs_fuzzy <- pFuzzyMatch(x = x,
                                y = y,
                                by.x = "orenames_x",
                                by.y = "orenames_y")

#Create synthetic data
x_orenames <- c("apple","oracle","enron inc.,"mcdonalds corporation")
y_orenames <- c("apple corp","oracle inc","enron","mcdonalds co")
x <- data.frame("orenames_x"=x_orenames)
y <- data.frame("orenames_y"=y_orenames)
z <- data.frame("orenames_x"=x_orenames[1:2], "orenames_y"=y_orenames[1:2])
z_true <- data.frame("orenames_x"=x_orenames, "orenames_y"=y_orenames)

# Perform merge
linkedOrgs_fuzzy <- pFuzzyMatch(x = x,
                                y = y,
                                by.x = "orenames_x",
                                by.y = "orenames_y")
```

`url2dt`*url2dt*

Description

Downloads a .zip file from a URL as a data.table from a URL.

Usage

```
url2dt(url, target_extension = ".csv")
```

Arguments

`url` character string with the URL housing the data object.
`target_extension` (default = ".csv") character string describing the target extension of the file in the downloaded .zip folder.

Details

`url2dt` downloads a zipped .csv file and loads it into memory based on the input URL.

Value

`z` The downloaded data object from the URL.

Examples

```
# Example download  
my_dt <- url2dt(url="https://www.dropbox.com/s/iqf9ids77dckopf/Directory_LinkIt_bipartite_Embeddings.csv.zi
```

Index

BuildBackend, [1](#)
BuildTransfer, [2](#)

GetCalibratedDistThres, [3](#)

LinkOrgs, [4](#)

pDistMatch, [7](#)
pFuzzyMatch, [8](#)

url2dt, [11](#)